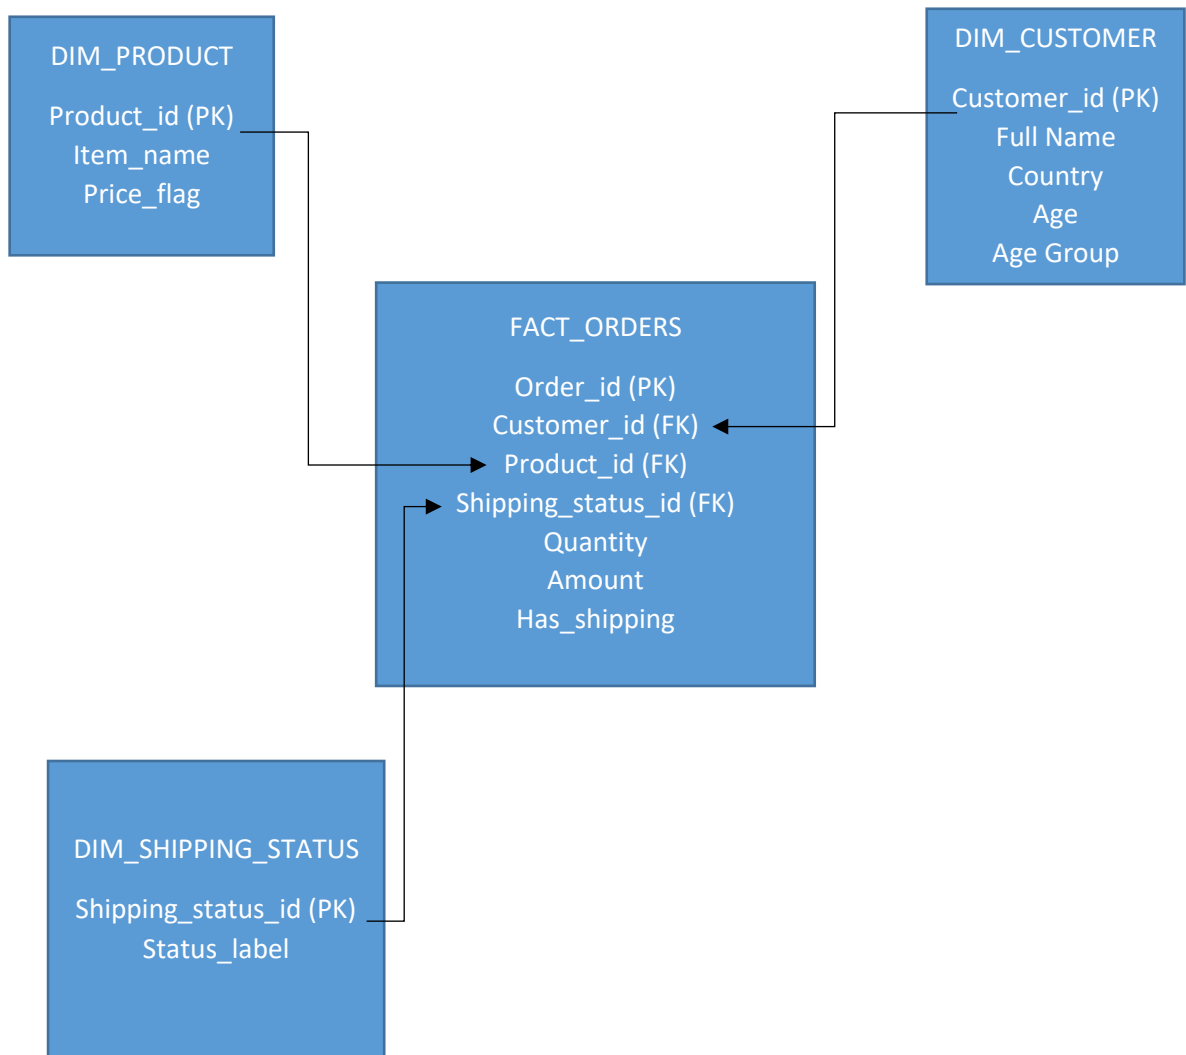


Model Design :



FACT_ORDERS

Field	Type	Description
Order_ID	PK	Unique identifier per order
Customer_ID	FK	Linked to DIM_CUSTOMER
Product_ID	FK	Linked to DIM_PRODUCT
Shipping_Status_ID	FK	Linked to DIM_SHIPPING_STATUS
Amount	Float	Amount paid
Quantity	Integer	Set as 1
Has_shipping	Boolean	True if there is matching record in shipping.json

DIM_CUSTOMER

Field	Type	Description
-------	------	-------------

Customer_ID	Integer	From source
Full_name	String	Concat(First,Last) from customer
Age	Integer	From source
Age_group	String	Derived : "<30", "30+"
Country	String	From source

DIM_PRODUCT

Field	Type	Description
Product_ID	Integer	System generated key
Item	String	Distinct item From Orders
Price_flag	Boolean	TRUE if multiple prices Exist

DIM_SHIPPING_STATUS

Field	Type	Description
Shipping_Status_ID	Integer	System generated key
Status_label	String	"Pending", "Delivered"

Story for Data Engineers:

Title : Build End-to-End Dimensional Model for Customer Order Data

As a Senior Data Analyst,

I want to ingest, cleanse, and model data from Customer.xls, Order.csv, and Shipping.json into a refined dimensional model consisting of 3 dimension tables and 1 fact table

so that business users can access high-quality, analysis-ready data that supports reporting on customer behavior, product sales, and delivery performance.

Tables :

1. DIM_CUSTOMER

- Filter invalid ages (< 0 or > 100)
- Remove NULL Customer_IDs
- Derive Full_Name = Concat (First, last) from customer
- Age_Group = case when age < 30 then '<30' else '30+' end

2. DIM_PRODUCT

- Deduplicate items
- Flag products with inconsistent pricing (e.g., Mousepad = 200, 250)

3. DIM_SHIPPING_STATUS

- Extract unique status values from Shipping.json

4. FACT_ORDER

- Join Order.csv to Customer.csv using Customer_ID
- Left join to Shipping.json (via Customer_ID)
- Flag Has_Shipping = FALSE where Shipping record is missing

Creation order:

- Create DIM_ tables first
- Create FACT_ORDER referencing ids
- Load in dependency order to ensure FK resolution

Load Frequency: One-time load for now (can be scheduled daily in production)

Add this timestamp as last field in each table.

Acceptance Criteria

- DIM_CUSTOMER, DIM_PRODUCT, DIM_SHIPPING_STATUS, and FACT_ORDERS tables are created with correct schema.
- All tables include Refresh_Timestamp
- Invalid or null Customer_ID, Age, and Country are excluded or flagged in DIM_CUSTOMER.
- Age_Group is derived correctly (<30, 30+).
- DIM_PRODUCT deduplicates items and sets Price_Flag = TRUE where multiple prices exist.
- Has_Shipping flag in FACT_ORDERS is set to FALSE if no matching shipping record.
- Referential integrity is maintained:
Customer_ID in FACT_ORDERS exists in DIM_CUSTOMER
Product_ID in FACT_ORDERS exists in DIM_PRODUCT
Shipping_Status_ID in FACT_ORDERS exists in DIM_SHIPPING_STATUS
- Orders without shipping are included (orphan shipping records are excluded).
- Row counts match expected values after cleaning.

Story for QA Team:

As a Data Analyst

I want to ensure that the dimensional model built by the Data Engineering team, including all fact and dimension tables, meets accuracy, completeness, and referential integrity standards **so that** the reporting and analytics layer is powered by reliable, trusted data that aligns with business logic and source definitions.

Tables to Validate

DIM_CUSTOMER
DIM_PRODUCT
DIM_SHIPPING_STATUS
FACT_ORDER

Test Scenarios:

Test	Table	Expected Results
keys are unique	All Dim tables	No duplicates in keys/ids
Columns not null	All	No nulls in critical fields
Data types match spec	All	Int, string, float etc. as defined
FACT_ORDER.Customer_id exists in DIM_CUSTOMER	Fact_order	Should be present
FACT_ORDER.Product_id exists in DIM_PRODUCT	Fact_order	Should be present
FACT_ORDER.Shipping_Status_ID exists in DIM_SHIPPING_STATUS	Fact_order	Should be present
Age group logic (<30, 30+)	DIM_CUSTOMER	Correctly labeled based on Age
Has_Shipping = TRUE only when shipping exists	FACT_ORDER	Derived correctly from joined data
Pricing conflicts flagged	DIM_PRODUCT	Products with multiple prices marked TRUE
Duplicate Order_IDs in source		Exclude or deduplicate
Orphan shipping records (no order)		Excluded or flagged
Orders with no shipping		Has_Shipping = FALSE
NULL Age or Country in DIM_CUSTOMER		Row excluded or marked invalid
Inconsistent Item name casing		Normalized in DIM_PRODUCT

Acceptance Criteria:

- No referential integrity breaks
- No critical NULLs
- Business rules consistently applied
- Edge cases (e.g., duplicate IDs, inconsistent prices) flagged or excluded
- Row counts, sum of financials between raw → refined match expectation